



NumaMMA

NUMA MeMory Analyzer

François Trahay – Télécom SudParis

Manuel Selva – INRIA

Lionel Morel – CEA

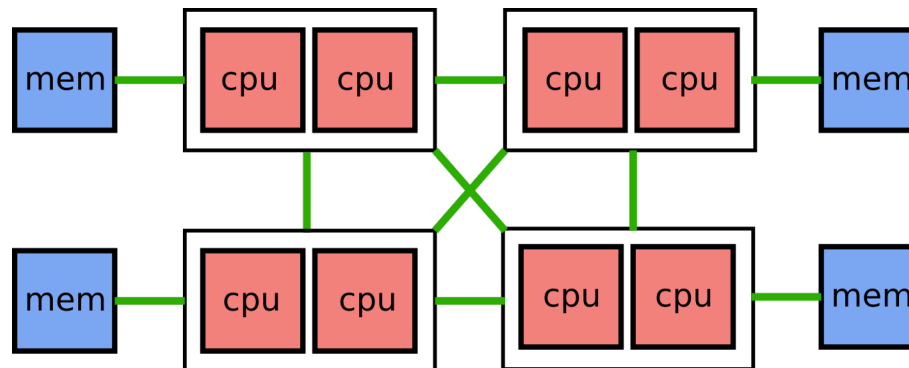
Kevin Marquet – INRIA



NUMA architecture

■ NUMA architectures are now common

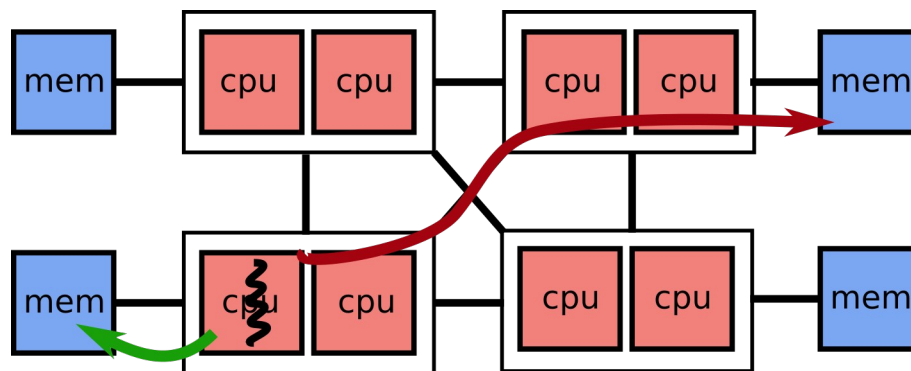
- Multi-socket systems
- Multicore CPUs
 - AMD Infinity Fabric (Zen CPU family)
 - Intel sub-NUMA clustering (Skylake family)
- Increase the available memory bandwidth



Impact of locality

■ NUMA = Non Uniform Memory Access

- Fast access to the local memory
- Slower access to remote NUMA nodes



→ **the locality of memory access impacts the performance**

eg. impact on NPB LU on a 48-core machine: up to 27%

→ **need to allocate pages on the right NUMA node**

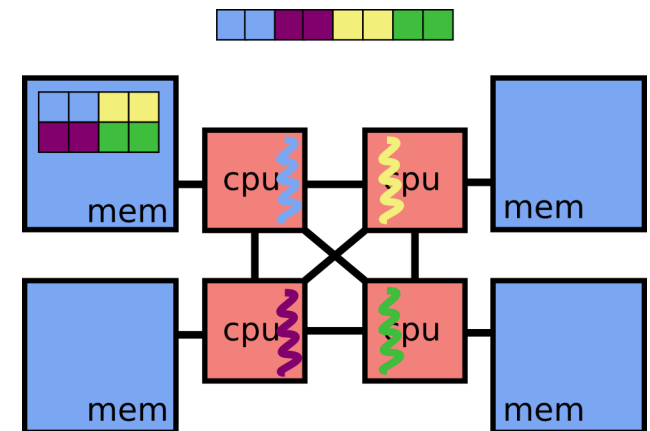
Memory allocation strategies

First-touch policy

■ *First-touch* policy

- Default policy on Linux
- Lazy allocation policy
- Allocate a page locally when a thread touches it
 - Assumption: this thread is probably the one that will use the page the most
 - Assumption may be wrong !

```
double *array = malloc(sizeof(double)*N);  
  
for(int i=0; i<N; i++) {  
    array[i] = something(i);  
}  
  
#pragma omp parallel for  
for(int i=0; i<N; i++) {  
    double value = array[i];  
    /* ... */  
}
```



Memory allocation strategies

Interleaved

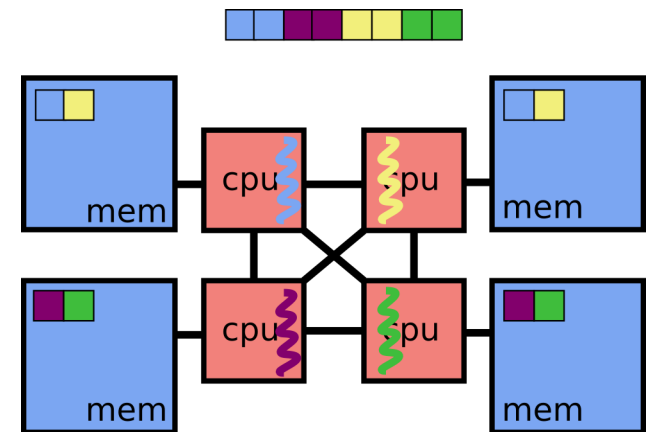
■ *Interleaved* policy

- The pages are allocated on multiple nodes in round-robin
→ Balance the load on multiple NUMA nodes

```
double *array = numa_alloc_interleaved(sizeof(double)*N);

for(int i=0; i<N; i++) {
    array[i] = something(i);
}

#pragma omp parallel for
for(int i=0; i<N; i++) {
    double value = array[i];
    /* ... */
}
```



Memory allocation strategies

Manual placement of memory pages

■ Manual placement of memory pages

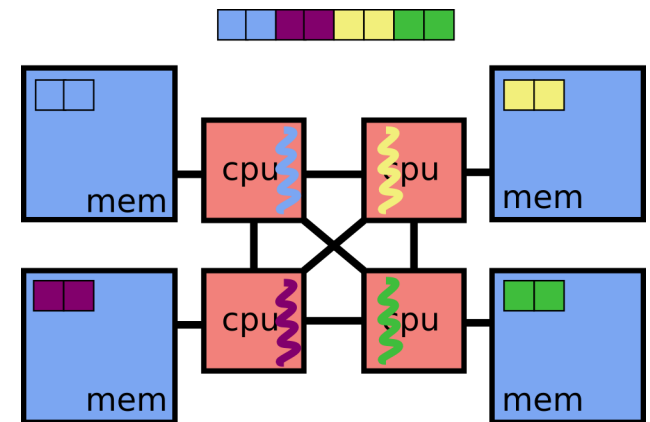
- Move pages to a specific NUMA node with `mbind()` or `move_pages()`
 - + precise control of page placement
 - manual placement of each page

```
double *array = malloc(sizeof(double)*N);

for(int i=0; i<N; i++) {
    array[i] = something(i);
}

mbind(&array[0], N/4*sizeof(double),
      MPOL_BIND, &nodemask, maxnode,
      MPOL_MF_MOVE);

#pragma omp parallel for
for(int i=0; i<N; i++) {
    double value = array[i];
    /* ... */
}
```



Choosing the best binding policy

- The threads access pattern may differ from one object to another

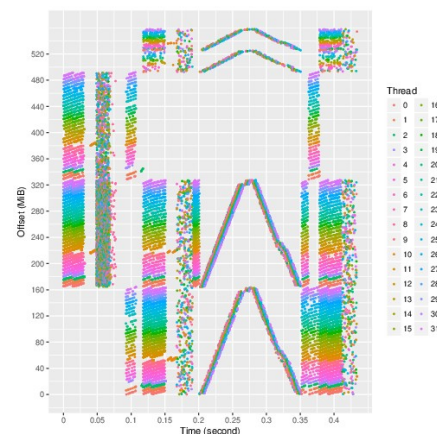
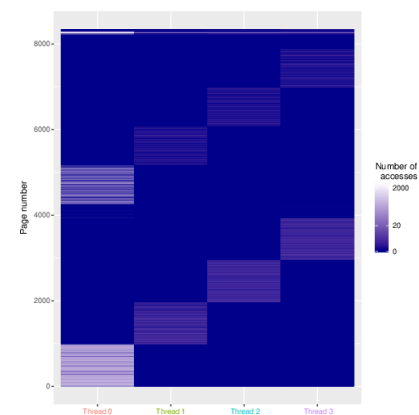
- eg. Matrix multiplication



- The best policy for one object may degrade the performance for another object
- Is it worth finding the best policy for one object ?
 - Some objects are rarely accessed

■ NumaMMA – NUMA MeMory Analyzer

- Analyze the memory access pattern of parallel applications
 - Low overhead collection of memory access
- Report:
 - The most accessed memory objects
 - Which thread access which part of an object
 - The evolution of access patterns over the time
- Freely available as open-source
 - <https://github.com/numamma>



■ Hardware memory sampling

- eg. Intel PEBS, AMD IBS
- Every X instructions, the CPU collects a sample
 - Address of the memory load/store
 - Thread/Instruction that issued the memory access
 - Where the data is stored (cache, RAM, remote RAM, ...)
 - Cost of the memory access (ie. latency)
- Information stored in a buffer
- Low overhead sampling (eg. < 1-2 %)

#tid	timestamp	address	mem_level	latency
0	5835423725112	0x557735cf07c8	L2 Hit	24
0	5835456302591	0x557736353ef8	Local RAM Hit	779
0	5835466068752	0x55773642a0c0	Local RAM Hit	657
0	5835471131886	0x5577362726e8	L2 Hit	23
0	5835566865010	0x557735d1fd28	L3 Hit	52
0	5835567586835	0x557735d04710	L3 Hit	64
0	5835604540592	0x557736320900	Local RAM Hit	1585
0	5835605025940	0x557735c39900	Local RAM Hit	265
0	5835618194705	0x557735f0e428	L2 Hit	24
0	5835693753719	0x557735f16a78	L2 Hit	23
0	5835709318658	0x557736260f00	Local RAM Hit	266

■ Static memory object

- eg. global variables
- Search for symbols in the ELF binary

■ Dynamic memory objects

- malloc, realloc, calloc, free, ...
- Intercept dynamic allocations with LD_PRELOAD

■ For each object, NumaMMA knows

- The allocation/de-allocation timestamp
- The start/end address

NumaMMA

Matching samples

- For each sample, find the matching memory object
- Once found, update counters
 - Number of read/write accesses
 - Total cost of memory accesses to the object
- For large objects, counters are computed per page
- Generate a summary of the most accessed objects

Summary of the call sites:

Sorting call sites

```
0  fields_ (size=2520000) - 34098 read access (total weight: 362881, avg weight: 10.642296). 66541 wr_access
1  [stack] (size=412316860415) - 47982 read access (total weight: 345827, avg weight: 7.207432). 60001 wr_access
2  constants_ (size=1272) - 589 read access (total weight: 5131, avg weight: 8.711375). 0 wr_access
3  /usr/lib/x86_64-linux-gnu/libgomp.so.1(+0x9b49) [0x7f6b06eb4b49] (size=192) - 96 read access (total weight: 672, avg weight: 7.0)
```

NumaMMA

Memory access patterns

■ Access pattern to an object

- X-axis: threads
- Y-axis: memory pages

■ Thread 0 access pattern

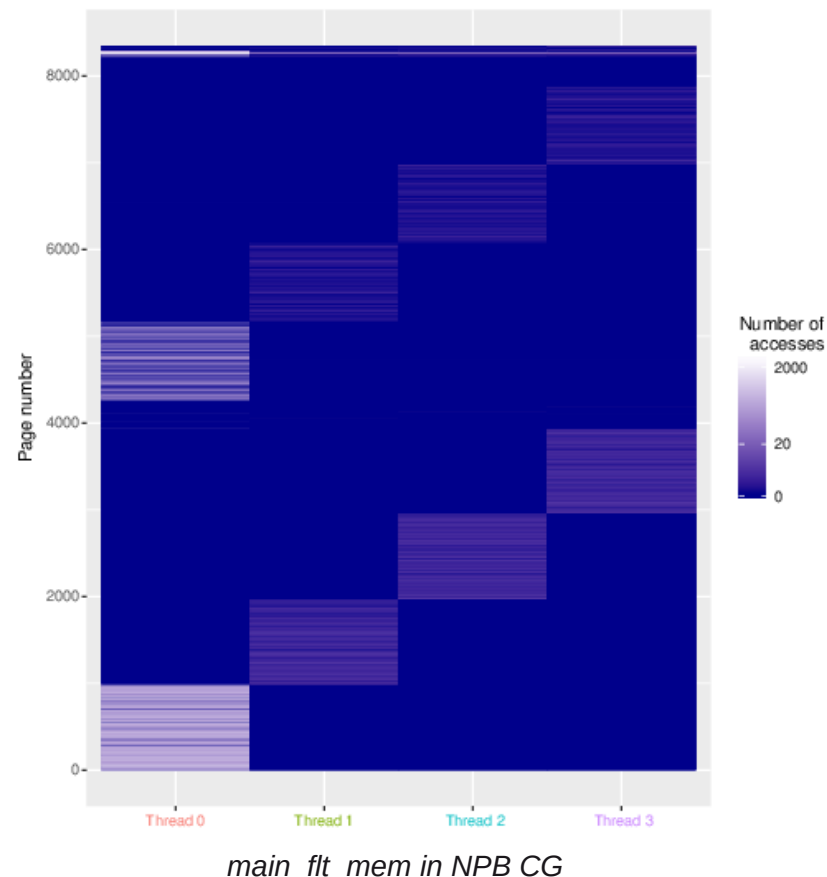
- Pages [0-992] [4267-5176]

■ Thread 3 access pattern

- Pages [993-1969] [5178-6074]

■ Due to sampling, some pages are not detected

- Depends on the sampling frequency



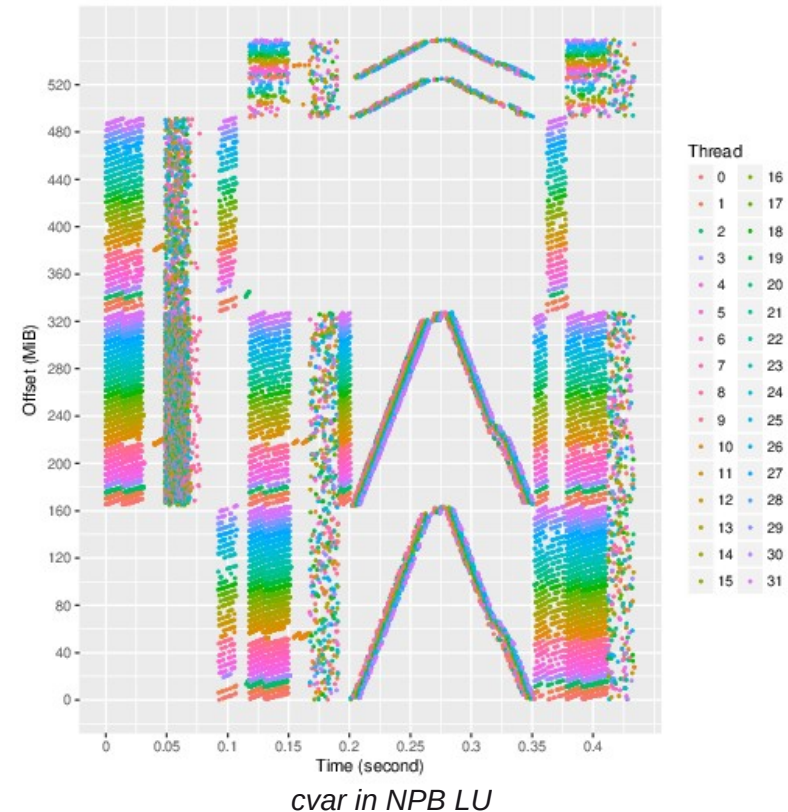
NumaMMA

Evolution of access patterns over the time

■ Access pattern to an object over the time

- X-axis: time
- Y-axis: memory pages
- Color: thread id

■ Allows to detect the phases of the application



Evaluation

■ Run multithreaded applications from NAS Parallel Benchmark and Parsec

- Evaluate the overhead
- How to use NumaMMA to improve the execution time of applications ?

■ Experiment setup

- Intel32:
 - 2 Intel Xeon E5-2630 v2 CPUs (16 cores/32 threads)
 - 32 GiB RAM (2 NUMA nodes)
 - Linux 4.11, GCC 6.3
- AMD48:
 - 4 AMD Opteron 6174 CPUs (48 cores)
 - 128 GiB RAM (8 NUMA nodes)
 - Linux 4.10, GCC 6.3

Evaluation Overhead

■ Overhead on NAS Parallel Benchmarks ■ 2 settings

- Running on Intel32
- OpenMP Implementation
- Threads are bound
- NumaMMA 2K
 - Sampling rate: 2000
 - Samples collected during malloc/free
- NumaMMA 10K
 - Sampling rate 10,000
 - Sample collection every 100ms

kernel	<i>native</i>	<i>NumaMMA_2k</i>			<i>NumaMMA_10k</i>		
	time(s)	time(s)	ovhd(%)	nsamples	time(s)	ovhd(%)	nsamples
BT.C	81.3	82.3	1.30	0.7 M	85.7	5.46	172 M
CG.C	21.6	23.3	8.01	0.7 M	22.2	2.93	41 M
EP.C	10	10.5	3.81	0.6 M	10.5	4.40	19 M
FT.C	19.6	20.2	3.14	0.5 M	21.7	10.77	39 M
IS.C	1.5	1.48	-2.11	0.18 M	1.45	-4.35	3 M
LU.C	61.8	58	-6.11	0.65 M	61.7	-0.12	110 M
MG.C	10.4	11.3	9.01	0.8 M	10.9	5.48	22 M
SP.C	168.6	169.8	0.68	0.5 M	169.6	0.59	334 M
UA.C	86.6	93.5	8.00	0.6 M	96.5	11.37	171 M

→ Low overhead

→ High resolution of a subset vs. Low resolution of the whole application

Evaluation

Case study: NPB LU

■ Application: NPB LU

- LU matrix factorization
- OpenMP Implementation

■ Most accessed objects:

<i>symbol</i>	<i>size</i>	<i>nb read</i>	<i>nb write</i>	<i>total</i>	<i>percent</i>
cvar	558 MB	112k	118k	230k	35.4
stack	?	105k	124k	229k	35.2
cexact	520 B	86k	0	86k	13.2
cjac	20 MB	39k	54k	94k	14.4
libgomp.so.1(+0x97c9)	8 KB	4k	97	4k	0.7

■ Access pattern analysis

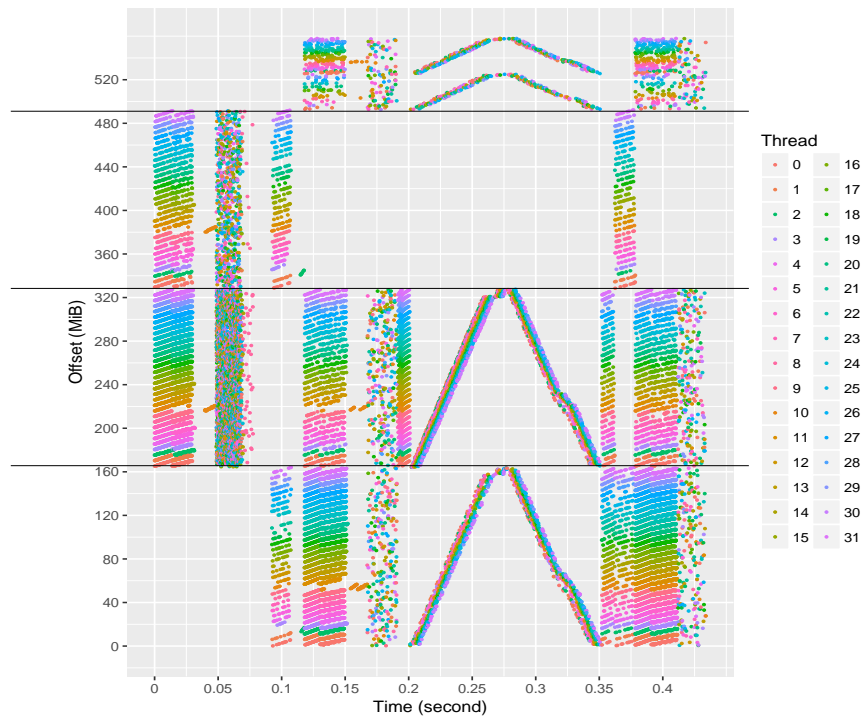
- `cexact` is accessed evenly by the threads → no easy optimization
- `cvar` and `cjac` could be optimized

Evaluation

NPB LU: Analyzing the access pattern

cvar

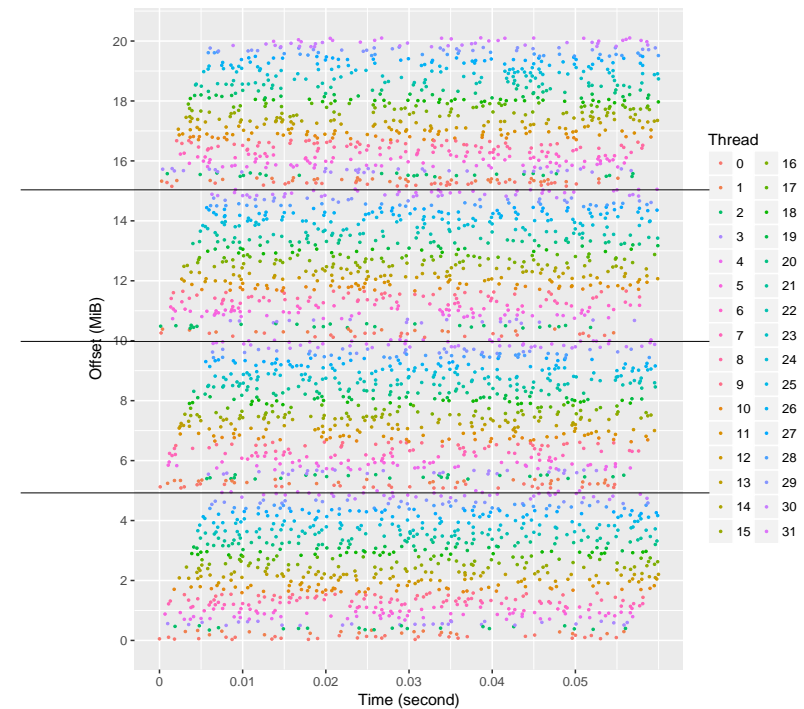
- 558MiB buffer
 - Accessed by slices of 160MiB
- **block-cyclic distribution**



Access pattern to cvar

cjac

- 20MiB buffer
 - Accessed by slices of 5MiB
- **block-cyclic distribution**



Access pattern to cjac

Evaluation

NPB LU: optimizing memory placement

■ Evaluation on AMD48

■ Comparing different memory placement

- First-touch: default policy
- Interleaved: interleave the memory pages of `cvar` and `cjac`
- Block-naive: use a block distribution for `cvar` and `cjac`
- NumaMMA: place `cvar` and `cjac` using a block-cyclic distribution

policy	execution time(s)	speedup
<i>first-touch</i>	102.53	1
<i>interleaved</i>	106.86	0.96
<i>block-naive</i>	109.88	0.93
<i>NumaMMA</i>	81.05	1.27

→ 27% performance improvement

Evaluation

Case study: Streamcluster

■ Application: Parsec Streamcluster

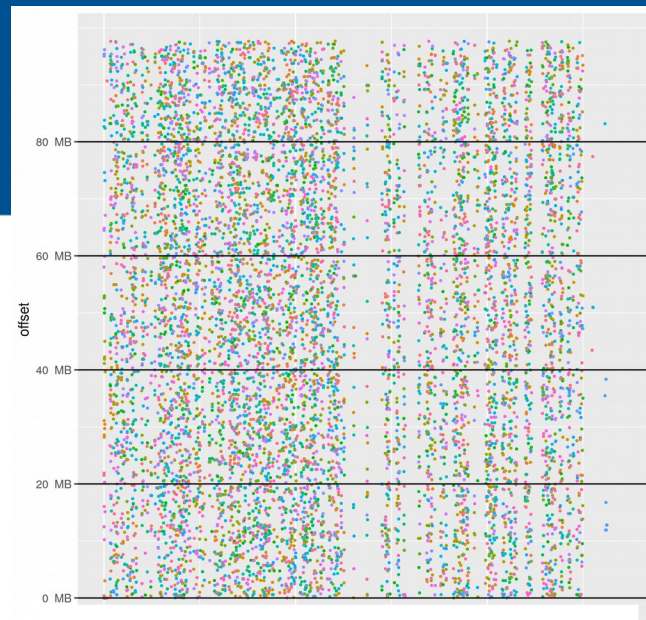
■ Most accessed objects

- `block` – 98 MiB buffer (66% of the samples)
 - Evenly distributed accesses
→ interleave pages
- `points` – 6 MiB buffer (31% of the samples)
 - Each thread accesses a part of the buffer
→ block distribution

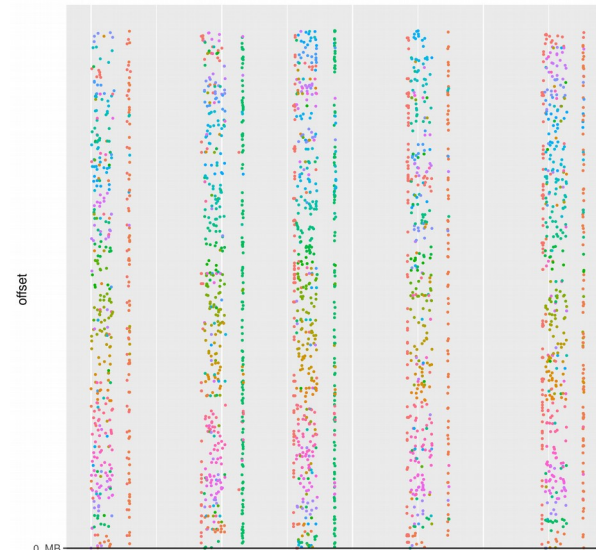
■ Evaluation

policy	execution time(s)	speedup
<i>first-touch</i>	93.72	1
<i>interleaved</i>	76.76	1.22
<i>block-naïve</i>	79.75	1.17
<i>NumaMMA</i>	73.32	1.28

- → 28 % improvement



Access pattern to block



Access pattern to points

Conclusion & future work

- **Memory placement is important for performance**
- **NumaMMA: NUMA MeMory Analyzer**
 - Use hardware sampling to collect memory access samples
 - Report the most accessed memory objects
 - Report the threads access pattern over the time
 - Available as open-source: <https://github.com/numamma/numamma>
- **Evaluation**
 - Low overhead (< 12%)
 - Reported information can be used for improving performance by up to 28%
- **Future work**
 - Port over AMD cpus
 - Use a signal to handle overflows
 - Automate memory placement